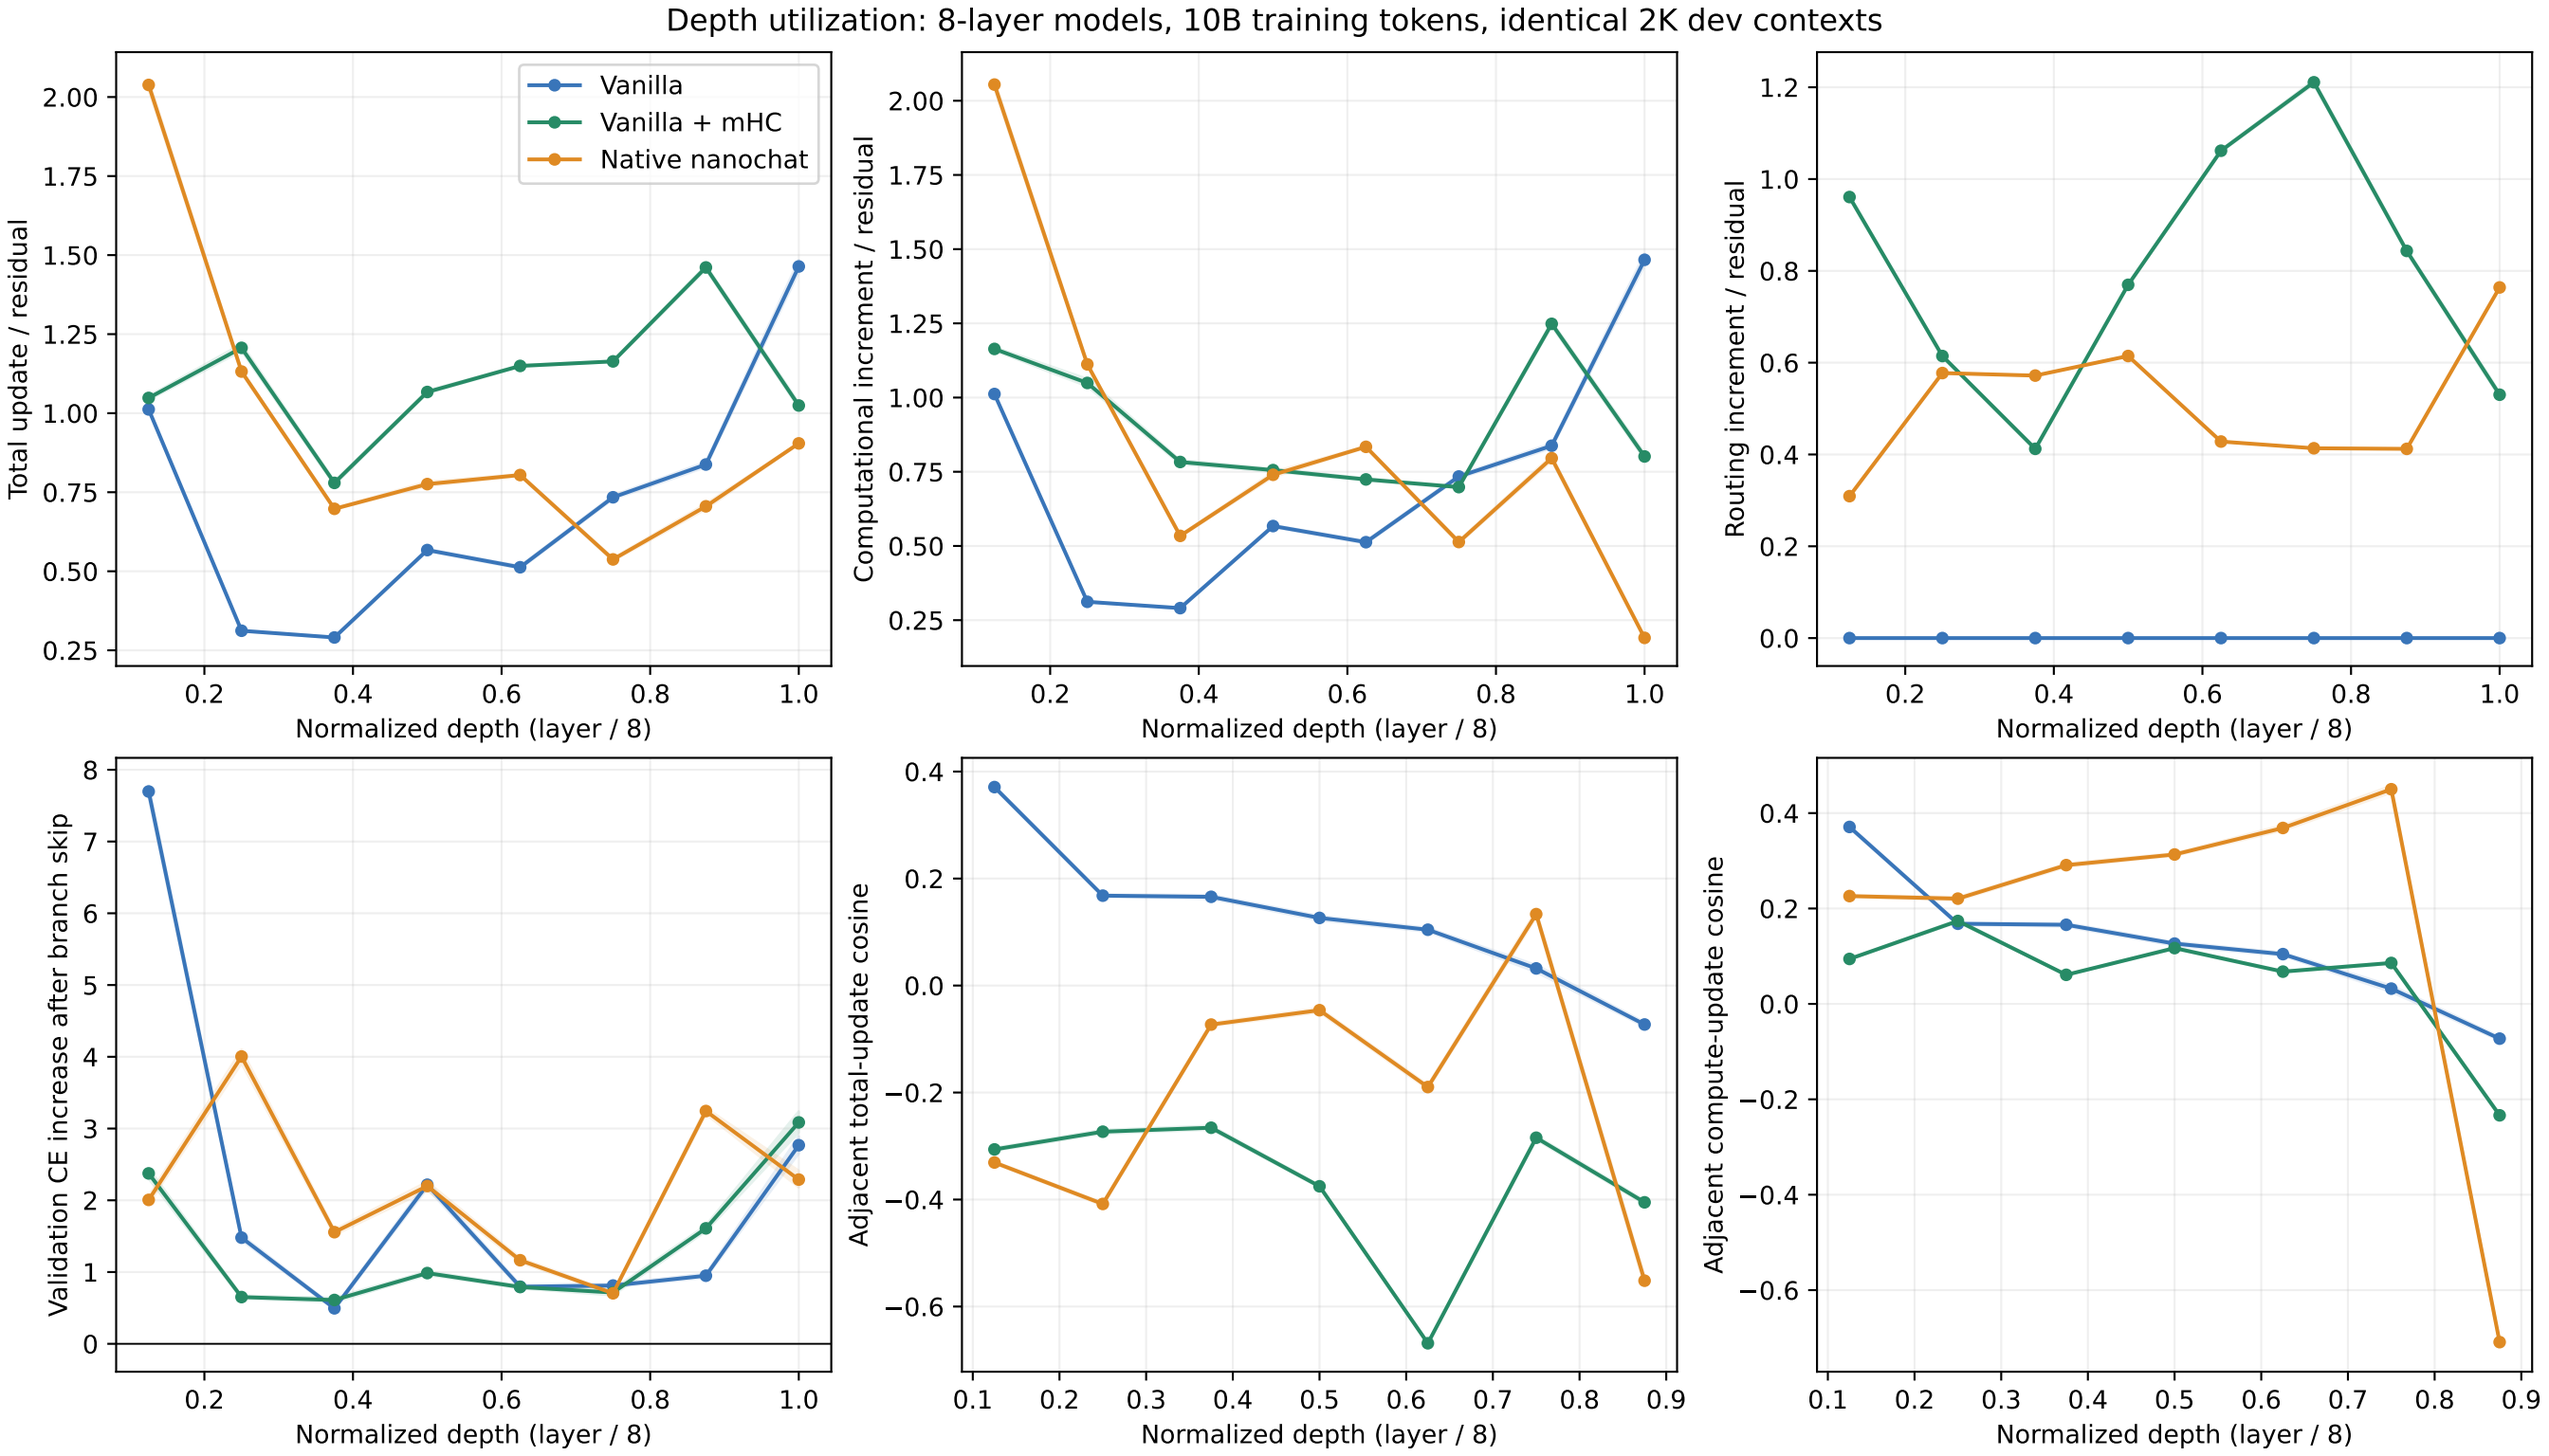
Depth utilization: 8-layer models, 10B training tokens, identical 2K dev contexts
Vanilla
Vanilla + mHC
Native nanochat
Total update / residual
Computational increment / residual
Routing increment / residual
Validation CE increase after branch skip
Adjacent total-update cosine
Adjacent compute-update cosine
Normalized depth (layer / 8)